

ZKLoRA: Efficient Zero-Knowledge Proofs for LoRA Verification



Bidhan Roy, Peter Potash, Marcos Villagra
Bagel Research Team*

January 18, 2025

Abstract

Low-Rank Adaptation (LoRA) is a widely adopted method for customizing large-scale language models. In distributed, untrusted training environments, an open source base model user may want to use LoRA weights created by an external contributor, leading to two requirements: (1) the base model user must confirm that the LoRA weights are effective when paired with the intended base model, and (2) the LoRA contributor must keep their proprietary weights private until compensation is assured.

We present **ZKLoRA**, a zero-knowledge verification protocol that relies on polynomial commitments, succinct proofs, and multi-party inference to verify LoRA–base model compatibility without exposing LoRA weights. ZKLoRA produces *deterministic* correctness guarantees and validates each LoRA module in **only 1–2 seconds** on state of the art large language models. This low-latency approach enables nearly real-time verification and promotes secure collaboration among geographically dispersed teams and contract-based training pipelines. The protocol ensures that the delivered LoRA module works as claimed, safeguarding the contributor’s intellectual property while providing the base model user with verification of compatibility and lineage.

1 Introduction

Large Language Models (LLMs) have attained remarkable success [1, 2], but verifying fine-tuned modifications such as LoRA [4] in an untrusted, distributed training environment can be difficult when the updated weights must remain private. Traditionally, one might re-run an entire forward pass or inspect thousands of parameters to ensure correctness, which is infeasible for massive models. ZKLoRA addresses this by generating a zero-knowledge proof of correctness for each LoRA module, guaranteeing that the private LoRA genuinely fits the base model. Crucially, **the verification stage for each LoRA module** in ZKLoRA remains about 1–2 seconds, even at scales of multi-billion parameter, state of the art large language base models.

¹Bagel is a research lab, using cryptography to make open source AI monetizable.

2 Preliminary Results

We benchmarked ZKLoRA across several LLMs and smaller models with different numbers of LoRA modules. Our central question is how verification times, as well as settings and proof generation times, grow with **the number of LoRA modules**, while also considering each LoRA’s average parameter size. Figure 1 and Table 1 detail this trade-off.¹




Figure 1: Total verification time (seconds) vs. number of LoRA modules, with dot size reflecting average LoRA size.

Base Model	# of LoRAs	Avg LoRA Size	Avg Settings	Avg Proof
distilgpt2	24	24576	38.0	31.6
openai-community/gpt2	48	49152	43.6	34.9
meta-llama/Llama-3.2-1B	32	26624	37.2	31.0
meta-llama/Llama-3.3-70B-Instruct	80	147456	54.9	46.9
meta-llama/Llama-3.1-8B-Instruct	32	163840	57.4	47.7
mistralai/Mixtral-8x7B-Instruct-v0.1	32	327680	86.1	73.7

Table 1: Model benchmark results for Settings and Proof generation time averaged by number of LoRA modules.

¹Note that the number of LoRA modules in a given model is not purely a function of the number of layers – it is also a choice of which weight matrices within each layer is targeted. For example, just targeting one matrix per layer (ie the Query matrix in the Attention Block) will yield one LoRA per layer. Conversely, targeting the Query, Key, and Value matrices yields three matrices per layer and makes the total number of LoRAs $3 \times num_layers$.

From Figure 1, we see that models with a higher LoRA count (e.g., 80 modules for a large 70B model) can indeed lead to larger total verification time overall. However, the slope remains modest due to ZKLoRA’s succinct design. For instance, even if each module is verified individually at around 1–2 seconds, verifying 80 modules can be done in just a few minutes, which is still practical for real-world usage.

Table 1 similarly shows that *proof generation* and *settings time* scale with the number of LoRA modules, but also depend heavily on each module’s average size (indicated by larger or smaller dot size). These two steps (proof generation on the LoRA contributor’s side and cryptographic circuit setup) can become more expensive, yet remain feasible in decentralized settings or paid contract relationships. The *Base Model User*, meanwhile, benefits from the relatively short verification overhead.

Overall, these plots confirm that ZKLoRA can handle multiple LoRA modules with minimal overhead for verifying correctness, emphasizing the viability of repeated or multi-adaptor scenarios in large-scale LLM pipelines.

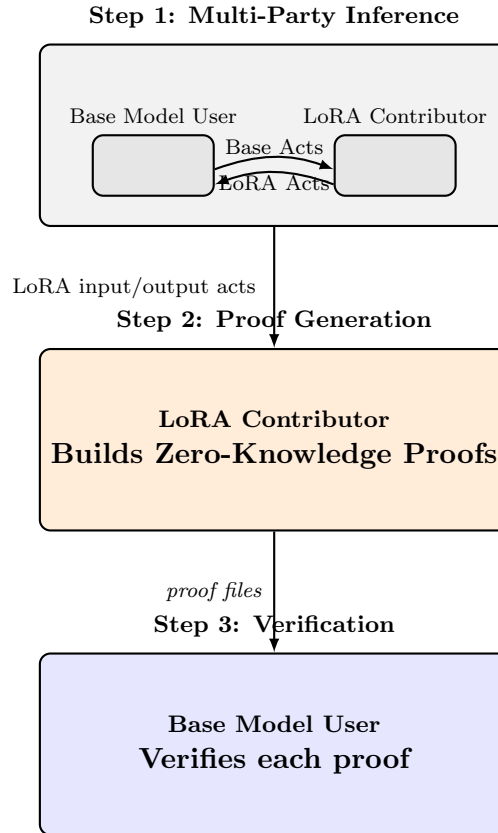


Figure 2: **Three-Step ZKLoRA Process (Vertical).** (1) Base Model User and LoRA Contributor exchange “Base Acts” and “LoRA Acts” in a full multi-party inference. (2) The LoRA Contributor generates cryptographic proofs for correctness. (3) The Base Model User verifies these proofs, ensuring correct LoRA alignment without revealing private adapter weights.




Figure 3: Flow in a Multi-Party Inference scenario between local base model and remote LoRA weights. The base model performs a local forward pass, computing **base_out** = $\mathbf{W}\mathbf{x}$. In parallel, the input \mathbf{x} is sent to the remote LoRA module, which returns $\Delta = \mathbf{B}\mathbf{A}\mathbf{x}$, where \mathbf{B}, \mathbf{A} are the low-rank finetuned matrices. The final output is **base_out** + Δ .

3 ZKLoRA

We implement a protocol that not only supports multi-party inference with partial activations exchanged between a base model user and a LoRA contributor, but also produces cryptographic proofs that the LoRA transforms are valid. First, the base model user processes any submodules that do not require LoRA directly on local hardware, sending partial activations only when a submodule contains LoRA-related layers. These exchanged activations, shown conceptually in Figure 2, correspond to “Base Acts” from the base model user and “LoRA Acts” from the LoRA contributor. Figure 3 also provides a fine-grained, computational perspective of what occurs when a submodule interfaces with remote LoRA weights.

Once the multi-party inference completes, the LoRA contributor shifts to a proof generation phase. At this stage, each LoRA module is compiled into a constraint system describing the LoRA transformations, and a key setup procedure yields the proving key, verification key, and possibly a structured reference string if the underlying zero-knowledge scheme requires one. The contributor then creates a “witness” by running partial activations through these constraints and finally produces the proof files. This entire flow is summarized in Algorithm 1, where Step 1 covers the multi-party inference, Step 2 addresses proof generation, and Step 3 handles the final verification check.

Once the proof generation is done, the base model user receives each proof and runs a fast verification procedure, typically requiring about 1–2 seconds per module. As Figure 3 suggests, this does not require the LoRA contributor to reveal the actual low-rank matrices. Instead, the contributor only sends updates and proofs that these updates conform to the declared LoRA transformations. If any single proof fails, the base model user can reject the entire LoRA submission; otherwise, the system is accepted as consistent with the underlying base model. This design reflects the synergy between LoRA’s parameter-efficiency and zero-knowledge cryptographic protocols: LoRA significantly shrinks the parameter footprint being proven, while the zero-knowledge aspect maintains confidentiality of the contributor’s proprietary weights. By merging these ideas, ZKLoRA enables trust-driven collaboration across decentralized infrastructures, contract-based training, and other scenarios where proof-of-correctness is essential but the LoRA code remains private. Our approach also builds on incremental verification concepts [7] and advanced proof systems such as Nova [6] and HyperNova [5], which allow us to scale proofs to large neural networks. Ultimately, this com-

ination provides a practical pipeline for parameter-efficient fine-tuning while verifying correctness in a succinct and minimally intrusive manner.

Algorithm 1 ZKLoRA Pseudocode

Require: *BaseModel* (public), *LoRAModel* (private), *Data*

Ensure: Verified outputs or **Reject**

```

1: Step 1: Multi-Party Inference
2: for each submodule  $s$  in BaseModel do
3:   if  $s$  contains LoRA layers then
4:     (a) run multi-party inference with LoRA Contributor for submodule  $s$ 
5:   else
6:     (b) run local inference on submodule  $s$  (no remote calls)
7:   end if
8: end for

9: Step 2: Proof Generation
10: for each LoRA module  $m$  in LoRAModel do
11:   (1) Circuit Compilation: parse LoRA-augmented layers, produce cryptographic circuit
12:   (2) Key Setup: generate settings, proving key, verification key, and SRS if needed
13:   (3) Witness Creation: run partial activations through circuit, record wire values
14:   (4) Proof: construct zero-knowledge proof  $\mathcal{P}_m$  for module  $m$ 
15: end for

16: Step 3: Verification
17: for each proof  $\mathcal{P}_m$  do
18:   if  $\text{Verify}(\mathcal{P}_m)$  fails then
19:     return Reject
20:   end if
21: end for
22: return Verified outputs

```

4 Related Work

4.1 Low-Rank Adaptation (LoRA)

Low-Rank Adaptation (LoRA) [4] is a technique for parameter-efficient fine-tuning of large language models (LLMs) that injects small, low-rank adapter matrices into specific layers of a pre-trained model. By isolating the fine-tuning process to these low-rank components, LoRA drastically reduces memory overhead compared to full-model fine-tuning. This design choice is especially appealing for massive LLMs where training or even storing all parameters can be prohibitive [3].

Beyond the clear advantage of reduced storage, LoRA also facilitates swapping multiple domain-specific adapters into a single base model, making it straightforward to maintain separate skill sets without instantiating an entire new copy of the model. These adapters can target specialized tasks (e.g., medical or legal text) with minimal overhead, driving LoRA’s widespread adoption. Yet verifying that a proprietary LoRA truly aligns with a base model (without revealing the adapter) remains problematic—precisely the gap ZKLoRA fills.

4.2 Incrementally Verifiable Computation

In a decentralized world, trust is a resource that is hard to achieve. In decentralized computation, we need to make sure the computations are both done and done correctly. In a seminal paper by Valiant (2008) [7], it was shown that proofs of knowledge can be used to assert the correct execution of general computations. That is, if M is a machine that runs for t steps producing a sequence of configurations c_0, c_1, \dots, c_t , then there exist an efficient and effective way to produce a computationally sound proof for the computation $c_0 \xrightarrow{t} c_t$. This idea is referred to as Incrementally Verifiable Computation or IVC.

The main goal of IVC is to produce compact, updatable proofs of correctness for a sequence of computations, so that each new step can be verified on its own while building on the guarantees of previous steps. This technique significantly reduces the verification overhead for long or evolving computations, which is invaluable in scenarios like decentralized networks, outsourced computation, and any application requiring frequent correctness checks.

Kumar et al. (2021) [6] introduced the proof system NOVA and the idea of recursive proofs, which are proofs that can “prove the correctness of other proofs.” Recursive proof composition is key to IVC where each proof attests to the correctness of both a step’s output and the validity of the previous step’s proof.

HyperNova [5] is a novel recursive argument system optimized for customizable constraint systems (CCS) that generalizes and improves upon prior approaches like Nova. It achieves efficiency through a folding scheme where the prover’s cryptographic costs are minimized and achieves zero-knowledge without relying on zkSNARKs. An IVC system allows the construction of proofs in zero-knowledge where the proofs reveal no information about the underlying computation or its inputs beyond the validity of the claim [7].

5 Conclusion

ZKLoRA provides a fast, robust mechanism to ensure that private LoRA modules remain consistent with a large base model. Our evaluations indicate that each LoRA module’s correctness can be verified in less than 2 seconds, even for multi-billion-parameter LLMs. This efficiency bridges the gap between privacy-preserving LoRA development and practical, real-time validation in large-scale deployments. In terms of future work, the most relevant and immediate work would be adding polynomial commitments of the base model’s activations (those that are sent as input to the LoRA contributor). This would take us one step closer to providing end-to-end verifiability of inference computation for LoRA-finetuned models. Other avenues of expansion could be integrating multi-contributor LoRAs, advanced zero-knowledge proofs for further performance gains, or partial data-privacy frameworks to shield user inputs as well as LoRA parameters.

References

- [1] Tom B Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. Language models are few-shot learners, 2020.
- [2] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding, 2018.

- [3] Ning Ding, Zhuosheng Zheng, Fei Tan, Yuxian Chen, Xipeng Xie, Zhiyang Liu, Xinze Dai, and et al. Delta tuning: A comprehensive study of parameter efficient methods for pre-trained language models, 2022.
- [4] Edward Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, and Weizhu Chen. Lora: Low-rank adaptation of large language models, 2021.
- [5] Rahul Kothapalli et al. Hypernova: High-degree polynomial folding via customizable constraint systems, 2024.
- [6] Abhiram Kumar, Mary Maller, Pratyush Mishra, Shashank Siri, Giovanni Tessaro, Pratik Vasudevan, and Yupeng Zhao. Nova: Recursive zero-knowledge arguments from folding schemes, 2022.
- [7] Leslie G Valiant. Incrementally verifiable computation or ivc. <https://dash.harvard.edu/handle/1/5026950>, 2008. Harvard University, Technical Report.